

TASK-5

INNER JOIN: returns only the rows where there is a match in both tables based on the specified condition. It excludes all rows that do not have corresponding entries in both joined tables. This is the most commonly used join when the goal is to fetch data that has a direct relationship between two tables. For example, joining customers with their orders will only show those who have actually placed orders.

Query		Query History	
1	SELECT	o.order_id, c.first_name, c.last_name, o.total_amount	
2	FROM	orders o INNER JOIN customers c	
3	ON	o.customer_id = c.customer_id;	

Data Output		Messages		Notifications	
order_id	first_name	last_name	total_amount		
integer	character varying (50)	character varying (50)	numeric (10,2)		
1	John	Smith	799.99		
2	Michael	Williams	59.99		
3	John	Smith	79.97		
4	John	Smith	39.98		

1	SELECT	c.customer_id, c.first_name, o.order_id	
2	FROM	customers c	
3	LEFT JOIN	orders o	
4	ON	c.customer_id = o.customer_id;	

Data Output		Messages		Notifications	
customer_id	first_name	order_id			
integer	character varying (50)	integer			
1	John	1			
2	Michael	3			
3	John	4			
4	John	5			
5	Robert	[null]			
6	Sarah	[null]			

LEFT JOIN: or LEFT OUTER JOIN returns all records from the left table and the matching records from the right table. If no match exists, the result will still include the row from the left table but fill the right side with NULLs. This join is useful when you want to retain all records from the primary (left) table, even if they don't have associated data in the right table.

Query		Query History	
1	SELECT	p.product_name, oi.quantity	
2	FROM	products p	
3	RIGHT JOIN	order_items oi	
4	ON	p.product_id = oi.product_id;	

Data Output		Messages		Notifications	
product_name	quantity				
character varying (100)	integer				
Smartphone X	1				
Blender	1				
Men's T-Shirt	2				
Women's Jeans	1				
Men's T-Shirt	2				

Query		Query History	
1	SELECT	c.first_name, o.order_id	
2	FROM	customers c	
3	FULL OUTER JOIN	orders o	
4	ON	c.customer_id = o.customer_id;	

Data Output		Messages		Notifications	
first_name	order_id				
character varying (50)	integer				
John	1				
Michael	3				
John	4				
John	5				
Robert	[null]				
Sarah	[null]				

RIGHT JOIN: is the reverse of a LEFT JOIN. It returns all records from the right table and the matched records from the left table. Unmatched rows from the right table will appear with NULLs for the left table's columns. This join is helpful when you prioritize data from the right table and still want to see unmatched entries from it.

FULL OUTER JOIN: combines the results of both LEFT and RIGHT joins. It returns all rows from both tables, filling in NULLs where there is no match on either side. This type of join is ideal when you want a complete dataset from both sources, including unmatched records from either table.

CROSS JOIN: produces the Cartesian product of two tables, meaning it pairs every row from the first table with every row from the second. No join condition is used. While it can be useful in generating test data or creating combinations, it can result in very large datasets and should be used with caution.

SELF JOIN: is when a table is joined with itself. This is useful for comparing rows within the same table, such as identifying hierarchical relationships (e.g., employees and their managers) or finding duplicates. It requires aliasing the table to distinguish between the two instances in the query.

Query

Query History

1

2

3

4

SELECT

p.p.product_name,

c.category_name

FROM

products p

CROSS JOIN

categories c

LIMIT

10;

Data Output

Messages

Notifications

≡+

SQL

	product_name character varying (100)	category_name character varying (50)
1	Men's T-Shirt	Electronics
2	Women's Jeans	Electronics
3	Blender	Electronics
4	Cookbook	Electronics
5	Smartphone X	Electronics
6	Wireless Earbuds	Electronics
7	Smart Watch	Electronics
8	Wireless Charger	Electronics
9	Men's T-Shirt	Clothing
10	Women's Jeans	Clothing

Query Query History

1

SELECT a.first_name AS customer1, b.first_name

2

AS customer2, a.city

3

FROM customers a

4

JOIN customers b ON a.city = b.city

Data Output Messages Notifications

Multi-table Joins: involve joining more than two tables in a single query. This is common in normalized databases where related data is spread across several tables. Each join condition connects one table to another, forming a chain. Proper join logic is essential to ensure accuracy and performance.

Query

Query History

1

▼

SELECT o.order_id, c.first_name, p.product_name, oi.quantity

2

FROM orders o

3

JOIN customers c ON o.customer_id = c.customer_id

4

JOIN order_items oi ON o.order_id = oi.order_id

5

JOIN products p ON oi.product_id = p.product_id;

Data Output

Messages

Notifications

≡

+

SQL

▼

📄

🗑

📧

📁

📥

📏

🔍

SQL

	order_id integer 🔒	first_name character varying (50) 🔒	product_name character varying (100) 🔒	quantity integer 🔒
1	1	John	Smartphone X	1
2	3	Michael	Blender	1
3	4	John	Men's T-Shirt	2
4	4	John	Women's Jeans	1
5	5	John	Men's T-Shirt	2

Joins with Filtering: apply conditions (using WHERE or ON) to refine the results of a join. This can limit the number of rows returned, focus on specific criteria (e.g., date range, status), and ensure data relevance. Filters can be applied before or after the join, depending on whether the condition is join-specific or general.

Query

Query History

1

▼

SELECT

p.product_name, p.price

2

FROM

products p

3

JOIN

categories c ON p.category_id = c.category_id

4

WHERE

c.category_name = 'Electronics';

Data Output

Messages

Notifications

≡+

▼

▼

SQL

	product_name character varying (100)	price numeric (10,2)
1	Smartphone X	967.99
2	Wireless Earbuds	181.49
3	Smart Watch	241.99
4	Wireless Charger	36.29

Query

Query History

1

▼

SELECT c.category_name,

2

COUNT(p.product_id) AS product_count

3

FROM categories c

4

LEFT JOIN products p

5

ON c.category_id = p.category_id

6

GROUP BY c.category_name;

Data Output

Messages

Notifications

≡

📄

▼

📋

▼

🗑

🗄

⬇

📈

SQL

	category_name character varying (50) 🔒	product_count bigint 🔒
1	Home & Kitchen	1
2	Electronics	4
3	Clothing	2
4	Books	1

Joins with Aggregation: combine rows across tables and then apply aggregate functions such as SUM, COUNT, or AVG. This is common in reporting and analytics, for instance, joining orders with customers and calculating the total spend per customer. Grouping is typically done using GROUP BY after the join.

Query		Query History	
1	▼	SELECT c.first_name, o.order_id	
2		FROM customers c	
3		JOIN (SELECT order_id, customer_id	
4		FROM orders	
5		WHERE total_amount > 100) o	
6		ON c.customer_id = o.customer_id;	
Data Output		Messages	Notifications
<div> <div>≡+</div> <div>📄</div> <div>▼</div> <div>📋</div> <div>▼</div> <div>🗑️</div> <div>🗄️</div> <div>⬇️</div> <div>📈</div> <div>SQL</div> </div>			
		first_name character varying (50) 🔒	order_id integer 🔒
1		John	1

Relationship Mapping: refers to the way entities (tables) are linked via keys, such as primary and foreign keys. Joins rely on this mapping to connect data logically, for example, mapping orders to customers using customer_id. Correct mapping ensures accurate data retrieval across relational databases.

NULL Handling in Outer Joins: In outer joins, NULLs are used to represent missing or unmatched values from one side of the join. It's important to account for these NULLs in logic, especially in filters and calculations. Functions like COALESCE are commonly used to handle or replace NULLs to maintain data integrity and avoid misinterpretation.

Join Optimization: focuses on improving the performance of join queries. Techniques include indexing join keys, reducing dataset size using filters early, selecting only necessary columns, and choosing the appropriate type of join. Poorly optimized joins can lead to slow queries and unnecessary resource consumption, especially with large datasets.